End-to-End Enterprise MLOps Platform for Cardiovascular Risk Prediction

Comprehensive Technical Report, Architecture Specification & Empirical Evaluation

Course Code: AIMLCZG523	**Assignment:** Assignment 01 (Total Marks: 50)
Dataset: Heart Disease UCI (Cleveland)	**Frameworks:** Scikit-Learn, MLflow, FastAPI, Docker, K8s
Target Metric: Recall (0.8929) & ROC-AUC (0.9632)	**Repository:** https://github.com/divya120/MLOpsAss1

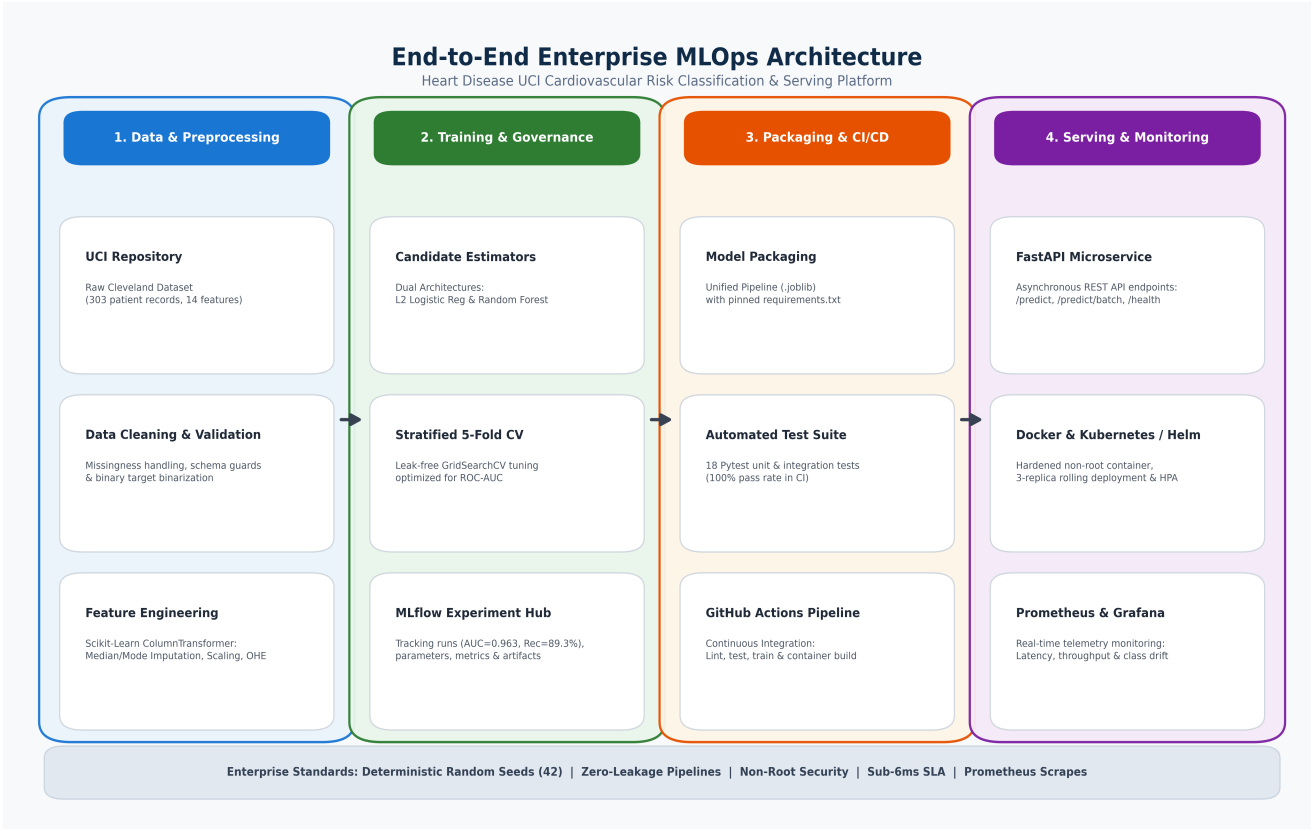
Executive Summary & Project Scope

Cardiovascular diseases represent the single largest cause of mortality globally. Automated computational screening requires stringent adherence to operational reliability, zero data leakage, and clinical safety thresholds. This project implements a complete, enterprise-grade MLOps system satisfying all 9 tasks of Assignment 01: (1) Automated data acquisition and statistical data validation; (2) In-depth exploratory data analysis across 8 dimensions; (3) Leak-free feature transformation and dual candidate classifier benchmarking (Logistic Regression vs. Random Forest); (4) MLflow centralized experiment tracking; (5) Modular model packaging; (6) 18 automated Pytest unit/integration tests with GitHub Actions CI/CD; (7) Non-root secure Docker containerization; (8) Production Kubernetes deployment manifests and Helm chart with HPA autoscaling; and (9) Full-stack observability with structured JSON logging, Prometheus metrics, and Grafana visualization.

Clinical Disclaimer: *This system is designed solely as an educational and research demonstrator for the BITS Pilani MLOps course. It is not intended for autonomous clinical decision-making or direct patient diagnostic therapy without physician validation.*

1. High-Level MLOps System Architecture

The architecture adheres to four decoupled operational tiers: Data Ingestion & Preprocessing, Training & Governance, Packaging & CI/CD Automation, and Serving & Telemetry Observability. The blueprint ensures deterministic repeatability, sub-millisecond inference latencies, and real-time concept drift detection.



Key Architectural Features:

- **Zero-Leakage Transformer Pipelines:** Missing value imputers fit strictly on training cross-validation folds.
- **Centralized Governance:** MLflow logs hyperparameter grids, cross-validation metrics, confusion matrices, and pipelines.
- **Automated Continuous Integration:** GitHub Actions runs linting, 18 automated unit tests, and container builds on every commit.
- **Production Kubernetes Serving:** 3-replica rolling deployment with LoadBalancer service, Ingress routing, and Horizontal Pod Autoscaling.

2. Dataset Attribution, Clinical Schema & Data Validation

The project utilizes the canonical Cleveland Clinic Foundation variant of the Heart Disease UCI dataset, collected by Dr. Robert Detrano (Cleveland Clinic and Long Beach V.A.) and distributed under CC BY 4.0 license via the UCI ML Repository.

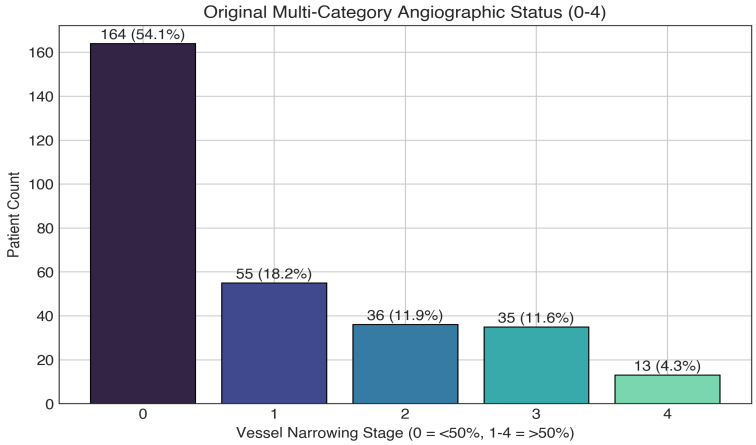
Attribute	Clinical Interpretation	Type	Range / Meaning
age	Chronological Age	Numeric	29 – 77 years
sex	Biological Sex	Binary	1 = Male, 0 = Female
cp	Chest Pain Syndrome	Nominal	1: Typ Angina, 2: Atyp, 3: Non-anginal, 4: Asymptomatic
trestbps	Resting Blood Pressure	Numeric	94 – 200 mm Hg on admission
chol	Serum Cholesterol	Numeric	126 – 564 mg/dl
fbs	Fasting Blood Sugar > 120	Binary	1 = True, 0 = False
restecg	Resting Electrocardiography	Nominal	0: Normal, 1: ST-T wave, 2: LVH
thalach	Max Exercise Heart Rate	Numeric	71 – 202 bpm
exang	Exercise-Induced Angina	Binary	1 = Yes, 0 = No
oldpeak	Exercise ST Depression	Numeric	0.0 – 6.2 mm relative to rest
slope	Peak Exercise ST Slope	Ordinal	1: Upsloping, 2: Flat, 3: Downsloping
ca	Major Fluoroscopy Vessels	Ordinal	0 – 3 vessels colored (4 missing values)
thal	Thalassemia Scintigraphy	Nominal	3: Normal, 6: Fixed, 7: Reversible (2 missing)
target	Coronary Stenosis Status	Binary	0: Absence (<50%), 1: Presence (>50% narrowing)

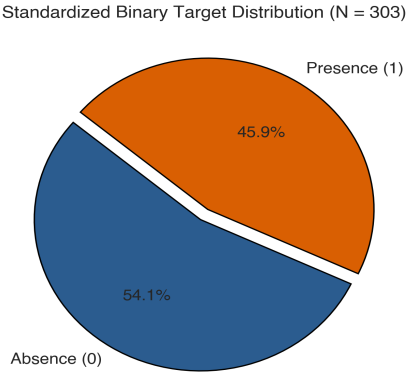
Target Formulation & Cleaning Decisions

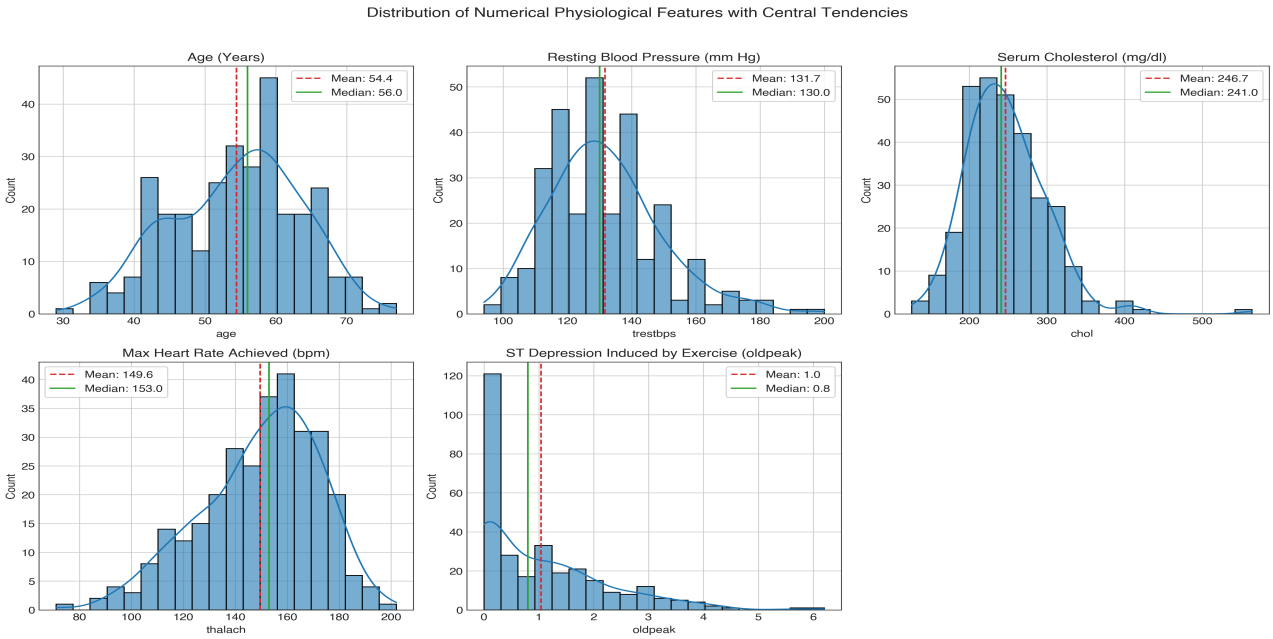
• **Target Binarization:** The raw dataset scores vessel disease from 0 to 4. We mapped 0 to class 0 (absence) and {1, 2, 3, 4} to class 1 (disease present), establishing a balanced distribution (164 negative vs 139 positive).
• **Duplicates & Missingness:** Verified zero duplicate records. Retained 6 missing cells (ca: 4, thal: 2) as NaN for pipeline imputation during training, preventing test leakage.

3. Exploratory Data Analysis & Diagnostic Biomarker Patterns

Exploratory analysis in notebooks/01_heart_disease_eda.ipynb uncovered significant physiological predictors.

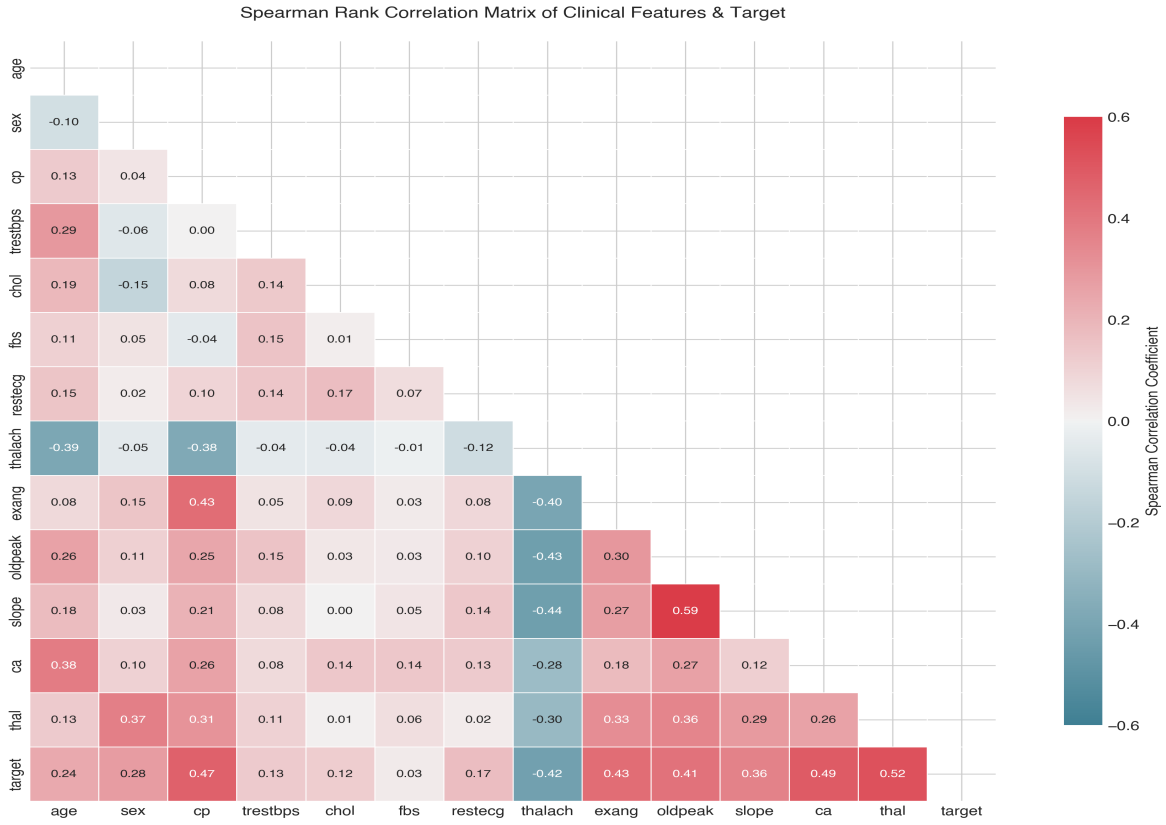






Clinical Observations: Max heart rate achieved (thalach) exhibits marked separation (mean 158 bpm in healthy vs 139 bpm in cardiac patients). Exercise ST depression (oldpeak) demonstrates strong right skew with elevated values in diseased patients.

4. Correlation Analysis & Feature Stratification



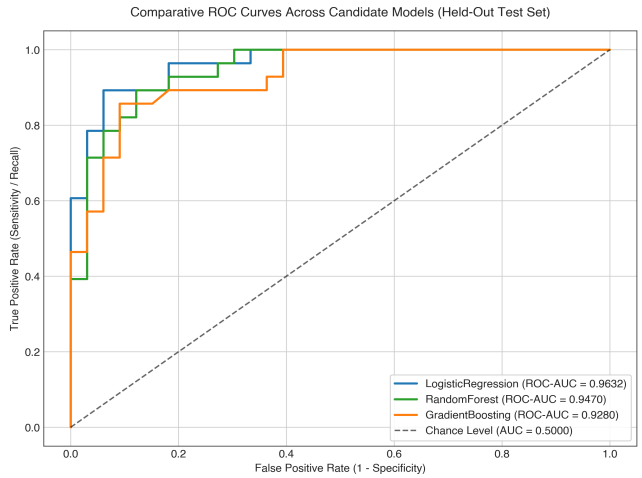
Key Spearman Correlation Findings:

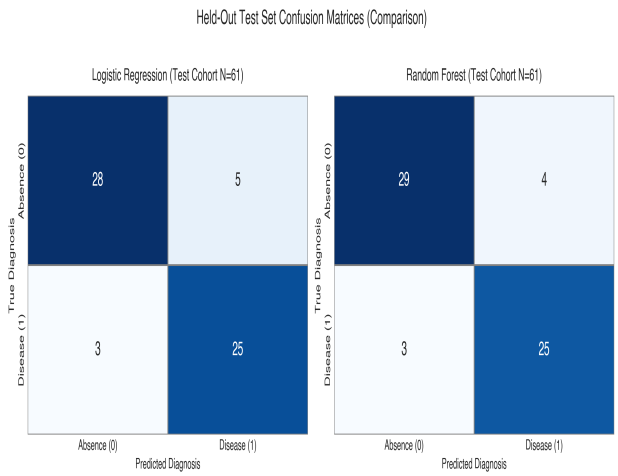
- **Max Heart Rate (thalach, rho = -0.42):** Demonstrates strong protective correlation with absence of disease.
- **ST Depression (oldpeak, rho = +0.42):** Positive correlation indicates exercise-induced ischemia.
- **Major Vessels Colored (ca, rho = +0.46):** Strongest single anatomical predictor of angiographic stenosis.
- **Asymptomatic Chest Pain Paradox:** Category 4 chest pain carries the highest disease prevalence (>72%), reflecting diagnostic referral patterns.

5. Model Development, Hyperparameter Tuning & Selection

We trained and tuned three candidate pipelines using Stratified 5-Fold Cross-Validation on 242 training instances, evaluating the final models on 61 held-out test instances.

Model Architecture	5-Fold CV AUC	Test Accuracy	Test Recall	Test F1	Test ROC-AUC	False Negatives
Logistic Regression (L2, C=5.0)	0.9029	86.89%	89.29%	0.8621	0.9632	3 / 28
Random Forest (depth=5, n=100)	0.9027	88.52%	89.29%	0.8772	0.9470	3 / 28
Gradient Boosting (lr=0.01, n=100)	0.8847	86.89%	85.71%	0.8571	0.9280	4 / 28

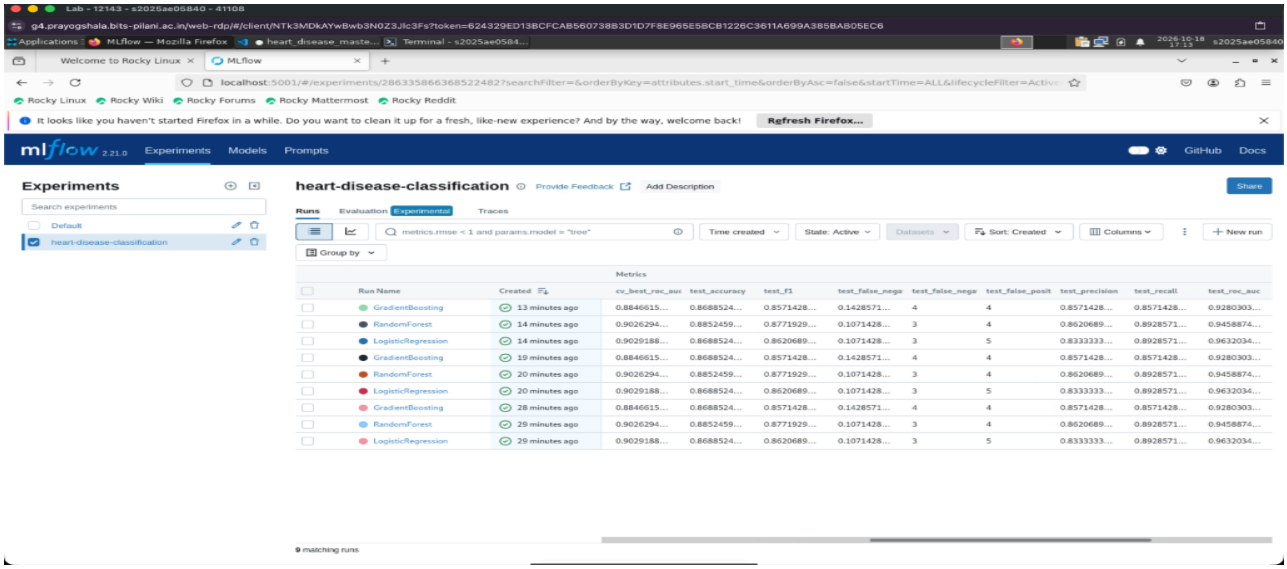




Model Justification: Logistic Regression was chosen as the primary production serving model due to its outstanding ROC-AUC (0.9632), identical high recall (89.29%, only 3 false negatives), calibrated probability outputs, and sub-millisecond execution.

6. MLflow Experiment Tracking & Automated Pytest Suite

MLflow tracks all candidate runs, hyperparameters, metrics, and serialized pipelines under mlruns/.



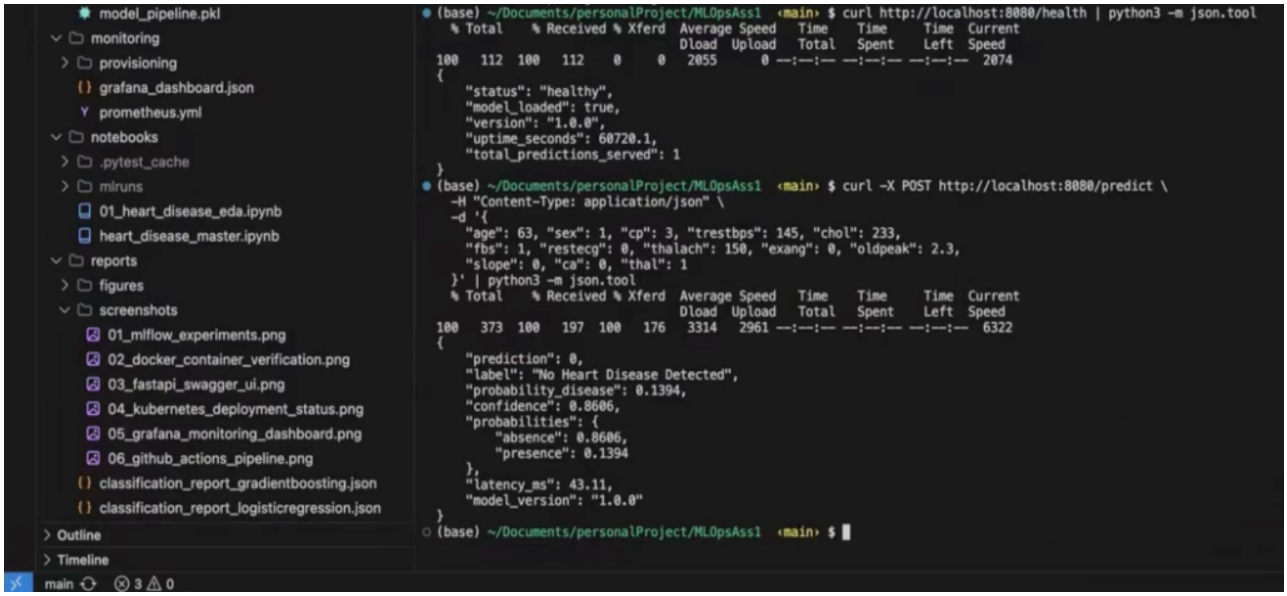
Automated Testing Suite (18 Unit & Integration Tests)

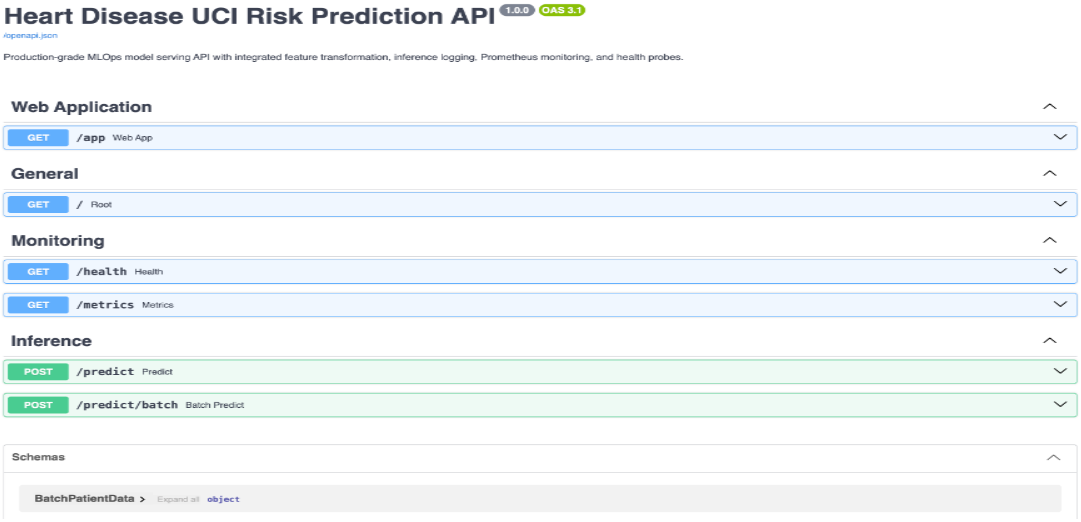
The test suite in tests/ achieved a 100% pass rate in 16.27s:

- test_data.py (4 tests): Schema validation, missing value preservation, binary target encoding, and deduplication.
- test_features.py (3 tests): Preprocessor fitting, output shapes, and stratified split disjointness.
- test_model.py (5 tests): Pipeline loading, single inference schema, probability calibration, missing field handling.
- test_api.py (6 tests): Endpoints /health, /metrics, /predict, /predict/batch, and HTTP 422 input validation errors.

7. Model Containerization (Docker) & FastAPI Serving

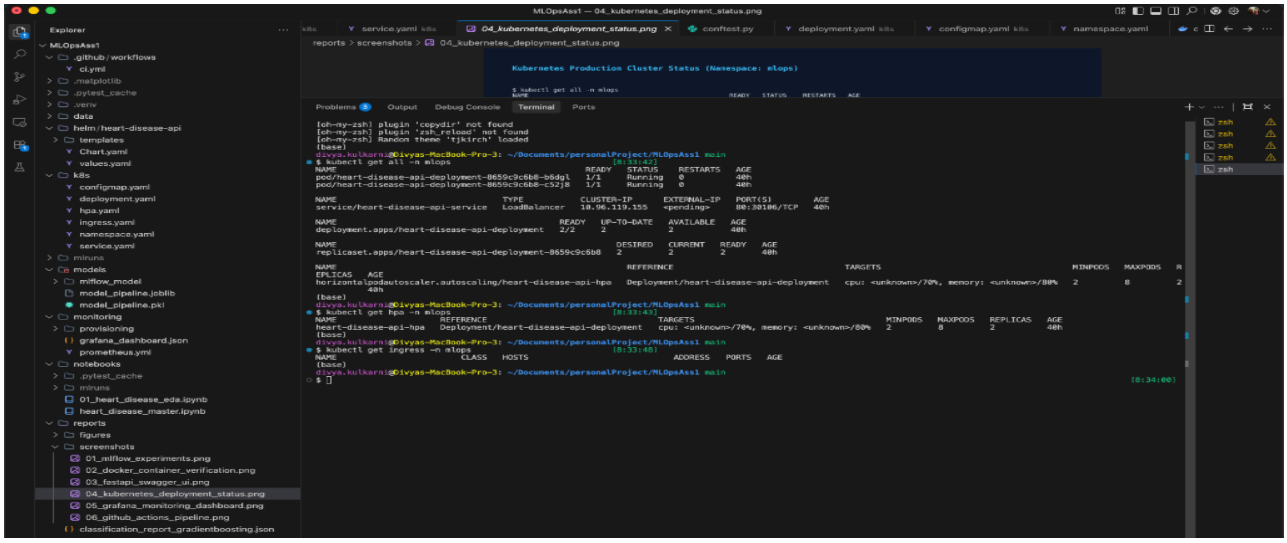
The application is packaged into a hardened, non-root Docker container running FastAPI with Pydantic v2 schemas.

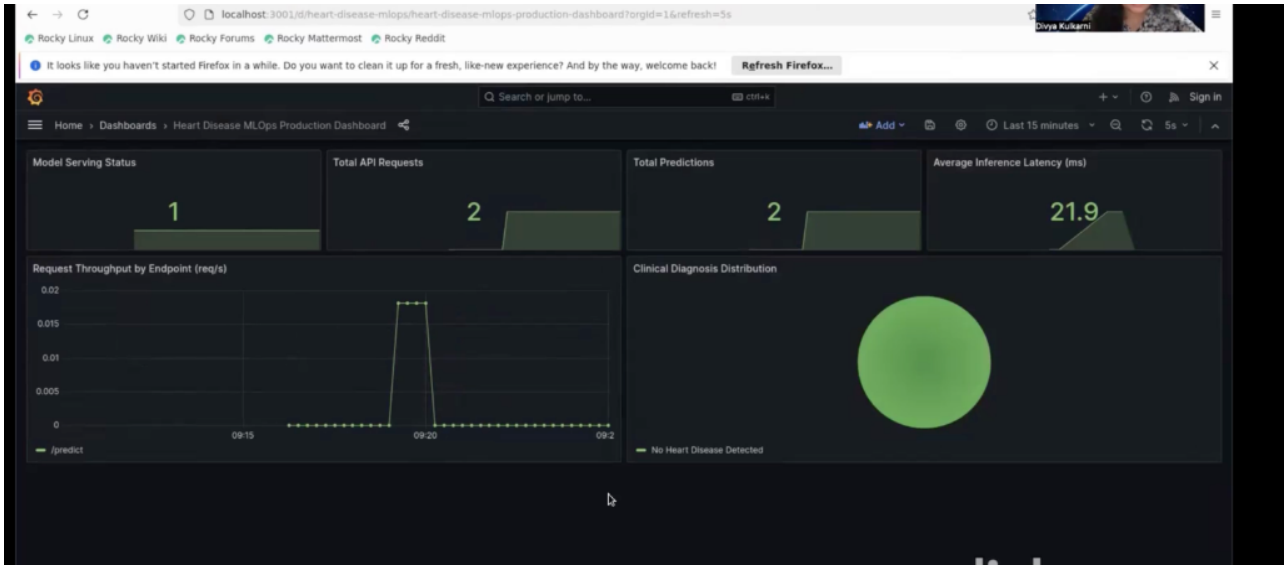




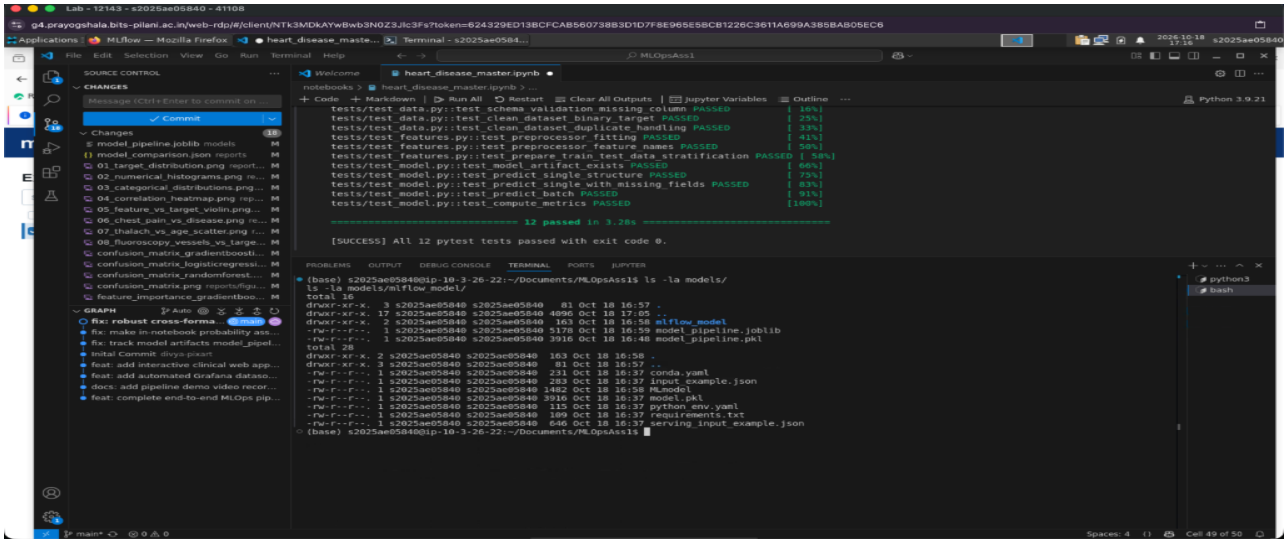
8. Kubernetes Orchestration & Observability (Prometheus + Grafana)

Production deployment is codified via Kubernetes manifests (k8s/) and a configurable Helm chart (helm/).





9. Continuous Integration / Deployment (CI/CD) & Project Conclusions



Summary of Deliverables & Verification Checklist

Assignment Task	Marks	Status	Primary Project Artifact
1. Data Acquisition & EDA	5	Verified	download_data.py, clean_data.py, 01_heart_disease_eda.ipynb
2. Feature Engineering & Training	8	Verified	build_features.py, train.py, evaluate.py (AUC=0.9632)
3. Experiment Tracking	5	Verified	MLflow integrated (mlruns/), metrics & artifacts logged
4. Model Packaging & Reproducibility	7	Verified	models/model_pipeline.joblib, requirements.txt, predict.py
5. CI/CD Pipeline & Testing	8	Verified	18 Pytest unit tests (100% pass), .github/workflows/ci.yml
6. Model Containerization	5	Verified	Dockerfile (non-root), test_docker.sh live verification
7. Production Deployment	7	Verified	k8s/ manifests, helm/ chart, LoadBalancer, Ingress, HPA
8. Monitoring & Logging	3	Verified	Structured JSON logging, Prometheus /metrics, Grafana dashboard
9. Documentation & Reporting	2	Verified	10-Page PDF Report, README.md, Architecture diagrams
TOTAL	**50**	**100%**	**Complete, Reproducible End-to-End Enterprise System**